

HemeWeb: Blood flow simulation in the cloud using docker

Steven Steven



Master of Science
School of Informatics
University of Edinburgh
2016

Acknowledgements

I would like to thank Indonesian Government, specifically Indonesian Endowment Fund for Education (*Lembaga Pengelola Dana Pendidikan / LPDP*) of Ministry of Finance. With their help, I can study in University of Edinburgh from start to end without any problems.

Declaration

I declare that this thesis was composed by myself, that the work contained herein is my own except where explicitly stated otherwise in the text, and that this work has not been submitted for any other degree or professional qualification except as specified.

(Steven Steven)

Table of Contents

1	Introduction	1
1.1	Motivation	1
1.2	Objectives	2
1.3	Outline	3
2	Background	4
2.1	Current HemeLB workflow	4
2.2	Cloud computing	5
2.3	Containerization Technology	5
2.4	Other High Performance Computing	5
3	Design and Specification	6
3.1	HemeWeb System Specification	6
3.2	HemeLB core container	7
4	Implementation	8
4.1	Deployment	8
4.2	HemeWeb web application	9
4.2.1	Running a simulation	9
4.2.2	Reproducing a simulation	11
4.2.3	Persistence	11
4.2.4	HemeWeb workflow	11
4.2.4.1	Pre-processing	11
4.2.4.2	Job configuration	11
4.2.4.3	HemeLB simulation	12
4.2.4.4	Post-processing	12
4.2.5	Challenge	12

List of Figures

2.1	Current HemeLB workflow taken from Steven [2016]	5
3.1	Planned HemeWeb architecture phase 1 from Steven [2016]	6
4.1	HemeWeb flow	11

Chapter 1

Introduction

Software are increasingly complex. Our everyday software are packed with features that makes its usage difficult. To people without familiarity with the product, it will be a barrier of entry to use it even when it helps them tremendously.

This also ties in to the complexity of a research that use these softwares. Open science dictates that research should be reproducible or replicable for it to better validate the research. However, recent findings have shown that not many research in psychology or even computation are replicable easily.

1.1 Motivation

To study how blood flow in a given vessel, Mazzeo and Coveney [2008] developed a fluid dynamic simulation software named HemeLB. Currently, it is actively developed and used by researchers to help their study. For example, Itani et al. [2015] used HemeLB for automated ensemble simulation of blood flow for a range of exercise intensities, Bernabeu et al. [2015] used it for detecting difference of retinal hemodynamics with regards to diabetic retinopathy, and recently Franco et al. [2015, 2016] used it to understand branching pattern of blood vessel networks.

As I have written in the proposal for this dissertation [Steven, 2016], HemeLB works by calculating fluid flow in parallel by using lattice-Boltzmann method [Mazzeo and Coveney, 2008]. This calculation allows HemeLB to simulate blood flow within a given blood vessel structure. Unfortunately, the calculation part is only a small part of the workflow to run the simulation. There are multiple pre-processing and post-processing steps needed to run the simulation from start to the end. These includes of preparing the input so HemeLB can work on it, and also processing the output so it is

ready to view.

These long pipelines of steps needed to run simulation, coupled with complexity of configuration of the software created a high barrier of entry for scientists and doctors to use HemeLB. Furthermore, as observed previously [Steven, 2016], interesting simulation will require parallel computing resources like ARCHER supercomputer which might be difficult to get access to by interested parties. While smaller simulation instances can run on typical laptop, most of the problems will require more powerful machines. These facts might prevent usage of the software by interested parties. More importantly, it shows there are still improvement that can be done to lower the barrier of entry for users to use HemeLB. This is important for HemeLB, especially when it is envisioned to be an integral part of future medical decision [Green, 2014].

Another aspect that HemeLB workflow can be improved is with regards to its reproducibility aspect. Researches that used HemeLB embrace reproducibility as one of its concern. As observed before [Steven, 2016], There are steps that are in place to make sure HemeLB and its simulation result are reproducible. First, the entire code base are publicly available on Github. Second, in running simulation with HemeLB, version of the software is automatically recorded. Lastly, in addition to the version used, input files and configurations are also recorded automatically. These facts can be seen from the publications mentioned above [Bernabeu et al., 2015; Itani et al., 2015; Franco et al., 2015, 2016] that include all these information. These information allow researchers interested to replicate the simulation to do it manually. Automation of these steps could further improve HemeLB's reproducibility and allow peers to replicate, duplicate, and audit its simulation results quickly and easily. This automation will be important, in addition to being usable, for HemeLB to become integral part of medical decision in the future.

1.2 Objectives

Based upon the needs to improve the usability, reproducibility, and auditability aspect of HemeLB project, I will develop a prototype web interface for HemeLB. This prototype web interface will lower barrier of entry in using HemeLB software compared to the current approach of using command line interface. In addition to that, using web interface will also allow features to be added to the simulation workflow which might not be essential to the HemeLB core itself. For example, automating packaging, sharing, and reproducing simulation result. These features are not essential for the

HemeLB core, but definitely help the overall workflow of blood flow simulation.

Using the dynamic capabilities of cloud computing vendor, the web backend should be able to dynamically scale without much efforts. On top of that, these infrastructures are available to everyone with a costs, allowing its user to access it without having to get access to supercomputers. Its user should be able to run a blood flow simulation without having to deal with the complexity of running each steps of the workflow manually.

In addition to the web interface, I will also develop deployment script so that peer could deploy its own instance of the web interface. This will ease up deployment process for individuals or organizations intending to use HemeLB for its own purposes. This script will be developed as part of the project.

1.3 Outline

I provide a brief introduction to the topic of this dissertation in this chapter. The rest of the chapters will be organized as follow:

- **Chapter 2.** I will provide background information that are necessary for readers to understand the concepts, technology and implementation that are done in this dissertation. HemeLB, containerization technology, cloud computing, High-Performance computing infrastructure, and other topics will be discussed in details in this chapter.
- **Chapter 3.** I will discuss the bulk of the work in this chapter. Implementation details and design of the proposed solutions will be provided and discussed in details.
- **Chapter 4.** Evaluation
- **Chapter 5.** Analysis
- **Chapter 6.** Future work

Chapter 2

Background

2.1 Current HemeLB workflow

Running a blood flow simulation using HemeLB currently consists of multiple steps. To understand how the proposed project can improve the current conditions, I will elaborate on how HemeLB workflow currently work.



Figure 2.1: Current HemeLB workflow taken from Steven [2016]

2.2 Cloud computing

2.3 Containerization Technology

2.4 Other High Performance Computing

Chapter 3

Design and Specification

I will discuss on the design and specification of HemeWeb. What are the resources needed to implement it and the design of the architecture.

3.1 HemeWeb System Specification

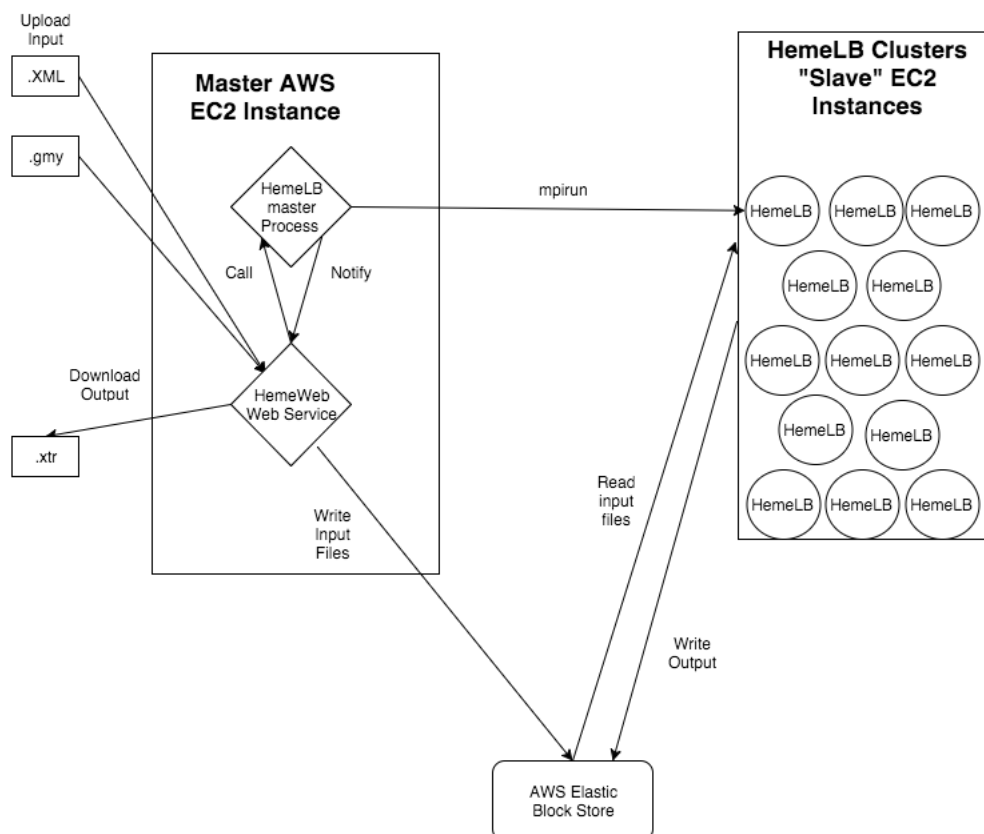


Figure 3.1: Planned HemeWeb architecture phase 1 from Steven [2016]

Figure 3.1 above was the original plan on how the system architecture will look like. The system will consist mainly of one master instance which will take care of starting up HemeLB cluster to compute the simulation. This master instance will run HemeWeb web application in which the user will interact with. Users will provide the web application with the inputs it needs to run. Which are the geometry file (.gmy) and the job configuration file (.xml).. With this input, the master instance will start up the appropriate HemeLB core cluster and run `mpirun` command to initiate the job execution. When execution is done, the output will be available for download to the users. In addition to that, the job files will be uploaded to amazon S3 so that it will be available for other instances of HemeWeb if it is shared, or for newly created instance with the same amazon AWS account.

3.2 HemeLB core container

Part of the designed system above is to have HemeLB cluster available for the HemeWeb master instance. In order to spin up the cluster easily, I decided to make a HemeLB core container only available in the docker hub (<https://hub.docker.com>). This container is based on an earlier work to ship HemeLB and all of its setup tools in a docker container so peers can run HemeLB without configuring their system a lot. However, for HemeWeb purpose, the whole suit of the setup tools are not essential for the simulation. What I planned to do is to strip it down to its bare minimum, so the container only contains the binary necessary for the simulation to run.

Chapter 4

Implementation

In this chapter, I will discuss HemeWeb's project implementation. This will consist on how the HemeLB core container is developed, how it is deployed, and how the web application itself is being implemented.

4.1 Deployment

In addition to have a working prototype available. I also wanted to make sure that the process of deploying the web application is as painless as possible. Therefore, I developed a deployment script for automatically configuring the architecture with minimal manual intervention. The script is made with the help of a configuration management and orchestration tools called ansible (<https://www.ansible.com>).

The goal of the script are as follow: 1. Provision the required master instance 2. Configure the master instances with the correct security and network settings 3. Configure and Install the HemeWeb correctly 4. Configure and install all the services needed by HemeWeb 5. Provision required compute instance 6. Configure the compute instances with the correct security and network settings 7. Configure the compute instances to run HemeLB core docker container

In the development of the script, I faced few problems in regard of orchestrating the infrastructure. Mainly with regards to cloud vendors having non-similar application programming interface to access or control the resources. This led to some part of the deployment needing manual intervention. One of this part is creating an image from the running instance so the master instance can request a new pre-configured instance instead of configuring from the start.

I developed the script as modular as I can so that it can be refactored easily. I

separated the cloud vendor specific instruction into its own folder for each vendor. This part mainly deal with the provisioning, security, and network configurations. The more common part, like installing softwares and services needed, are abstracted into the common scripts that are called by the vendor specific script after provisioning instances are done. With this in mind, I had been able to run the deployment script on three different cloud vendors. They are Google Cloud Platform, Amazon Web Service, and Digital Ocean. After the deployment script is done, the architecture will be ready to run HemeLB simulation on the cloud.

This script is online as part of the codebase that is published on github. It is available on <https://github.com/SeiryuZ/HemeWeb/tree/master/deployment>.

4.2 HemeWeb web application

This section will discuss the bulk of the work that is developed for the project. The main interface for users to interface with HemeWeb, the web application. I developed this web application using django web framework for my familiarity with the framework, so I can focus my time on developing the web application.

4.2.1 Running a simulation

The main feature of HemeWeb web application is the ability to run HemeLB simulation. Allowing users without technical know how to run the simulation on command line interface to use web browser to run it. The web application take two input files, geometry file and HemeLB configuration file and store it on the master interface.

The web interface will then allow users to modify the HemeLB configuration file with an in-browser text editor and also configure the job execution like compute instance count, compute instance type, and the HemeLB core container to use. After the configuration is done, the job will be queued into a queue system which are based on a redis Pub/Sub mechanism.

An asynchronous workers (Different from the web application workers) will then pick up the queued up job. The worker will run an ansible script to startup correct amount of compute unit from the cloud vendor. For this project, amazon web service is the only available cloud vendors. The compute unit will be started up from the state after the configuration on the deployment part, so it will not waste too much time to configure the base image. However, this is not yet ready to run the HemeLB simulation.

What the script will do next is to reconfigure the compute unit more, like reconfiguring docker service to point to the correct master address, mounting the remote file system containing the input files, pull the correct HemeLB core container from docker hub, and run the container.

After all the reconfiguration processes are done, the master's asynchronous worker will then fire an mpi job for HemeLB with the correct parameters. HemeLB simulation will run and produce output which will be written to the shared folder with the master instance. The compute instances will be terminated after the simulation is done.

Outputs of the simulations are then made available for the users to be downloaded via the web interface.

These features are part of the original scope of the web application as I planned on my proposal [Steven, 2016]. However, I extended the web application so that it can handle more cases. The first thing I added were the ability to handle the pre-processing of the input files. The input files that are used by the HemeLB simulation are generated from a geometry generation step that are done before the simulation. This geometry generation step took different input files, a geometry file (.stl) and a profile file (.pr2) to generate the input that HemeLB simulation can parse.

The way I implement the pre-processing stuff is to add another form for user to add input files to create new job. It receive the .stl and .pr2 file, save them, and queue up an asynchronous job that will pre-process these input into the correct files that will be feed into the HemeLB simulation configuration part.

Also, I added post-processing feature to the web application. The output files generated by the HemeLB simulation are effective to write in parallel. However, these files are not directly viewable by software like VTK viewer. These files need further post-processing, this is where the post-processing step is introduced in HemeWeb. I added the post-processing step, piping the outputted files into two python scripts, into the asynchronous worker. So after the HemeLB simulation step is done, it has an extra responsibility to run the post processing step on the master interface. The outputted files then will be packaged with the original output for download by the user.

4.2.2 Reproducing a simulation

4.2.3 Persistence

4.2.4 HemeWeb workflow

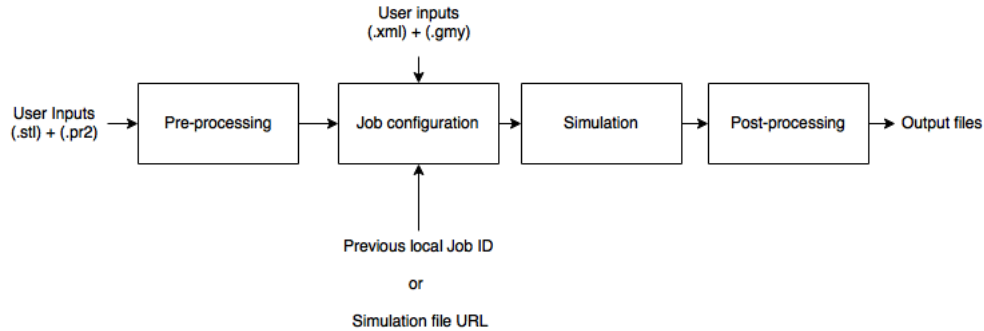


Figure 4.1: HemeWeb flow

Figure 4.1 illustrate how the HemeWeb web application works. it consists of mainly 4 core activity that will be discussed in details in the following section.

4.2.4.1 Pre-processing

HemeWeb handle pre-processing of inputs that are needed so that HemeLB simulation can parse the files. User provide a geometry file (.stl) and a profile file (.pr2) to the web application. HemeWeb will then create a job instance with these two files, save them locally on master instances and queue the pre-processing job.

The asynchronous worker on master instance will then pick up the job whenever they are free. It will run the pre-processing python script to generate the geometry files and HemeLB configuration file. These files will then be saved on the master instance, and HemeWeb will track these files by recording the path to these files on the job instance. Now the job instance is ready for the next step. of the workflow.

4.2.4.2 Job configuration

In this step, HemeWeb application will take a job instance with correctly set geometry file (.gmy) and HemeLB configuration (.xml). However, there are multiple ways that HemeWeb can get this correctly set job instance. As illustrated on Figure 4.1 , there are 4 possible entry points for this step. They are:

- **From the post-processing step.** These files are generated from the previous pre-processing step. The job instance is directly used in this step
- **User's provided geometry and configuration file.** User have pre-processed their own file locally, or have their own geometry and configuration files available. HemeWeb will create a new job instance, save both files and keep them tracked with the job instance.
- **User's provided previous job ID.** There are two possible case when user specify previous job ID. First, the previous job is available locally on the HemeWeb instance. Second, the previous job is cached on the persistent storage on the cloud vendor and are not available locally. HemeWeb will download the previous file from the persistent storage if it is not available locally. It will then create a new job instance that copy the previous job's geometry file and configuration file to be used for further configuration.
- **User's provided simulation file URL.** The last alternative is for user to provide the simulation file URL. Simulation files are uploaded to a persistent storage at the end of the workflow. These files, if made public, can be used by other instance of HemeWeb to download the simulation files and use it as a basis to create a new job instance. The way the system work is the same as using previous job ID, but its source is not its own persistent storage, but other people's simulation files.

After the job instance is created from one of the four way possible discussed above, HemeWeb will then ask users for the job configuration. This entails on configuring how many instances should be started, what are the type of the instance, and the hemelb-core container version. After configuring all these parameters, then the user can queue the job into the queue system.

4.2.4.3 HemeLB simulation

4.2.4.4 Post-processing

4.2.5 Challenge

One problem that I faced on developing the web application is that currently, it is developed mainly for the use on amazon web service (AWS). Due to the limited time for

the development, I haven't been able to develop the features on multiple cloud vendors. The features like, automatically uploading simulation file and restoring previous simulations, are tied in the code to the amazon infrastructure. It should be possible to abstract these features to handle multiple cloud vendors given more time.

Bibliography

- Bernabeu, M. O., Lu, Y., Lammer, J., Aiello, L. P., Coveney, P. V., and Sun, J. K. (2015). Characterization of parafoveal hemodynamics associated with diabetic retinopathy with adaptive optics scanning laser ophthalmoscopy and computational fluid dynamics. In *2015 37th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 8070–8073. IEEE.
- Franco, C. A., Jones, M. L., Bernabeu, M. O., Geudens, I., Mathivet, T., Rosa, A., Lopes, F. M., Lima, A. P., Ragab, A., Collins, R. T., et al. (2015). Dynamic endothelial cell rearrangements drive developmental vessel regression. *PLoS Biol*, 13(4):e1002125.
- Franco, C. A., Jones, M. L., Bernabeu, M. O., Vion, A.-C., Barbacena, P., Fan, J., Mathivet, T., Fonseca, C. G., Ragab, A., Yamaguchi, T. P., et al. (2016). Non-canonical wnt signalling modulates the endothelial shear stress flow sensor in vascular remodelling. *Elife*, 5:e07727.
- Green, C. (2014). Computer simulation could become 'integral' in the diagnosis of.
- Itani, M. A., Schiller, U. D., Schmieschek, S., Hetherington, J., Bernabeu, M. O., Chandrashekar, H., Robertson, F., Coveney, P. V., and Groen, D. (2015). An automated multiscale ensemble simulation approach for vascular blood flow. *Journal of Computational Science*, 9:150–155.
- Mazzeo, M. D. and Coveney, P. V. (2008). Hemelb: A high performance parallel lattice-boltzmann code for large scale fluid flow in complex geometries. *Computer Physics Communications*, 178(12):894–914.
- Steven, S. (2016). Hemeweb: Container based high performance computing scenario in cloud infrastructure for hemelb. Informatics research proposal, University of Edinburgh.